

SecureScan AI

Source Code Bug & Vulnerability Detection System

"Detect vulnerabilities before attackers do."

Document Type	Phase 1 — Project Roadmap & Software Requirements Specification (SRS)
Course	AI335L — Deep Learning Lab
Instructor	Sir Haseeb
Semester	Spring 2026
Submission Date	April 27, 2026
GitHub Repository	LINK

TEAM MEMBERS

#	Full Name	Registration ID
1	Member 1 — Ali Abbas Shah	S2024cs012
2	Salman Tanveer	S2024cs019
3	Hammad Ali	S2024cs021

SECTION 2 | EXECUTIVE SUMMARY / ABSTRACT

Section 2: Executive Summary

 **What Problem Are We Solving?**

Millions of developers worldwide now use AI coding assistants — tools that generate code instantly but completely ignore security. The result is production software riddled with hidden vulnerabilities that attackers can exploit. Manual code review is slow, expensive, and error-prone. Existing free tools flag issues but cannot explain them or suggest fixes.

 **What Is SecureScan AI?**

SecureScan AI is a deep learning system that reads Python or JavaScript source code and automatically detects bugs, security vulnerabilities, and coding errors — then explains each problem in plain English and suggests a corrected version of the code. A developer pastes their code into a web interface and receives a full security report within seconds.

Our approach combines three complementary neural network architectures:

- CodeBERT (Transformer) — A pre-trained model by Microsoft that already understands code the way humans do. We fine-tune it on vulnerability data.
- LSTM (Recurrent Network) — Reads code sequentially, line by line, capturing patterns that span multiple lines — just like an experienced code reviewer.
- MLP (Feedforward Network) — The final decision layer that classifies the bug type, vulnerability category, and severity (Low / Medium / High / Critical).

~215,000 Training Samples	≥ 75% Target F1-Score	3 DL Architectures	6 Weeks Delivery Timeline
-------------------------------------	---------------------------------	------------------------------	-------------------------------------

The system is trained on approximately 215,000 real-world vulnerable code samples from two leading research datasets — BigVul and Devign — and deployed as a free Gradio web application requiring no installation. All training runs on Kaggle's free GPU tier, making this project fully reproducible without paid cloud resources.

SECTION 3 | PROBLEM STATEMENT & MOTIVATION

Section 3: Problem Statement & Motivation

3.1 The Modern Software Security Crisis

The global software industry is facing an accelerating security crisis. According to the National Vulnerability Database (NVD), over 28,000 new software vulnerabilities were publicly disclosed in 2023 alone — nearly double the number from five years earlier. The rapid adoption of AI-powered coding assistants (GitHub Copilot, ChatGPT, Gemini) has dramatically increased code production speed, but these tools are optimized entirely for functionality, not security.

The financial and operational consequences are staggering:

- The average cost of a data breach in 2024 reached USD \$4.88 million (IBM Cost of a Data Breach Report, 2024).
- Software vulnerabilities account for over 60% of all publicly disclosed breaches.
- The global application security market is projected to reach USD \$25 billion by 2027, yet most individual developers and small teams cannot afford enterprise-grade tools.
- The 'time-to-detect' window for vulnerabilities in production systems averages 197 days — nearly seven months of exposure.

3.2 Who Is Affected?

Affected Group	Their Problem	Why Current Tools Fail Them
Individual Developers	Use Copilot / ChatGPT to write code fast — unaware of hidden security flaws	Linters only catch style/syntax. Security tools require expensive licenses.
CS Students & Learners	Learning to code but no security training. Ship vulnerable projects.	No affordable, beginner-friendly security checker exists.
Small Dev Teams	No dedicated security engineer. Ship products with unreviewed AI code.	Enterprise tools (Snyk, Veracode) cost \$500–\$2,000+/month.
Open Source Maintainers	Receive unreviewed contributions from unknown contributors worldwide.	Manual review doesn't scale. No automated explanation tool exists.

3.3 Where Existing Solutions Fall Short

- **Linters:** Only catch syntax errors and style violations. They cannot detect SQL injection, buffer overflows, or authentication bypass vulnerabilities. Think of them as spell-checkers — useful, but not security tools.
- **Enterprise SAST Tools:** Powerful but require complex setup, organization-wide licenses costing thousands of dollars per year, and dedicated security engineers to interpret results. Completely inaccessible for individual developers or students.
- **Free Scanners:** Can flag specific known patterns but do not use AI to understand code context, cannot explain why something is dangerous in plain language, and offer no fix suggestions.

- **AI Coding Assistants:** Generate code but explicitly disclaim security responsibility. They optimize for "does it run?" not "is it safe?"

The Gap We Fill

No free tool currently combines: (1) AI-based deep understanding of code context, (2) detection of both bugs and security vulnerabilities, (3) plain-language explanation of what the issue is and why it matters, and (4) an actionable corrected version of the code — all in a single, zero-installation web interface. SecureScan AI fills this gap for the millions of developers who cannot access or afford enterprise security tooling.

SECTION 4 | LITERATURE REVIEW

Section 4: Literature Review

This section surveys twelve works — six foundational studies and six current 2026 publications — that directly inform our system design. For each paper, we summarize the problem addressed, the technical approach, the dataset used, the reported results, and — critically — the limitation that justifies our work.

4.1 Paper Summaries

Paper 1: CodeBERT (Feng et al., 2020) — Foundation Paper

- **Problem Solved:** How can a single model understand both natural language (English) and programming languages (Python, Java, JavaScript, etc.) at the same time?
- **Approach:** A Transformer-based model pre-trained on 2.1 million code-documentation pairs using two training objectives: Masked Language Modeling (predicting hidden words/tokens) and Replaced Token Detection. Think of it as GPT but trained specifically on code.
- **Dataset:** CodeSearchNet — 2.1 million functions with paired natural language descriptions across 6 programming languages.
- **Results:** Achieved state-of-the-art performance on code search, documentation generation, and code summarization benchmarks. The model understands what code does, not just what it looks like.
- **Limitation for Our Work:** CodeBERT was NOT fine-tuned on vulnerability data. It understands code semantics but cannot distinguish safe from unsafe code patterns without further training on security datasets.

Paper 2: BigVul (Fan et al., 2020) — Dataset Paper

- **Problem Solved:** The lack of a large, realistic, labeled dataset for vulnerability detection research.
- **Approach:** Scraped GitHub commits that fixed CVE-reported vulnerabilities. Extracted the "before" (vulnerable) and "after" (fixed) versions of each code function. Each sample is labeled with the CVE ID, CWE type, and severity score.
- **Dataset:** BigVul itself — 188,000 C/C++ vulnerable code functions, the largest publicly available vulnerability dataset.
- **Results:** Became the standard benchmark dataset for vulnerability detection research. Widely cited and used in dozens of subsequent papers.
- **Limitation:** Almost entirely C/C++ code. Limited Python or JavaScript coverage, which are the languages most commonly used with modern AI coding assistants.

Paper 3: LineVul (Fu & Tantithamthavorn, 2022) — Most Relevant Recent Work

- **Problem Solved:** Not just detecting that a file has a vulnerability, but pinpointing the exact line number where the vulnerability exists.
- **Approach:** Fine-tuned CodeBERT + attention mechanism analysis. The attention weights from the Transformer model are used to highlight which specific tokens/lines contributed most to the vulnerability prediction.
- **Dataset:** BigVul — the full 188,000 sample dataset.
- **Results:** F1-score of approximately 91% on vulnerability detection; strong line-level localization performance — closest prior art to our project.

- **Limitation:** Single programming language (C/C++). Requires significant GPU compute for fine-tuning. Provides no fix suggestion, no plain-language explanation, and no deployment as a user-facing tool.

Paper 4: Deep Learning-Based Source Code Vulnerability Detection — Systematic Review (2026) — Current State of the Field

- **Problem Solved:** As LLMs accelerate code generation and propagation of vulnerabilities, existing surveys were outdated. This work systematically reviews deep learning for source code vulnerability detection over the last four years.
- **Approach:** Queried six major academic databases (IEEE Xplore, ACM DL, Google Scholar, arXiv, CNKI, Wanfang) using predefined keywords, screening ~6,600 publications down to 430 eligible studies through a four-phase filtering process.
- **Dataset:** Meta-analysis of 430 primary studies across multiple datasets including BigVul, Devign, NVD, and SARD.
- **Source:** Published in Springer Cybersecurity, February 2026. [\[Read paper\]](#)
- **Relevance to Our Work:** Directly validates our choice of CodeBERT + LSTM hybrid and the BigVul + Devign dataset combination as the current best practice for vulnerability detection. Confirms that multimodal deep learning approaches are the dominant 2026 direction.
- **Limitation:** Being a survey, it does not propose a new model — it maps the landscape. Our work operationalizes the top-recommended approach into a deployed tool.

Paper 5: Vul-RAG — Enhancing LLM-based Vulnerability Detection via Knowledge-level RAG (Du et al., 2026) — ACM TOSEM

- **Problem Solved:** LLMs struggle to distinguish vulnerable code from similar-looking patched code, achieving only 0.06–0.14 accuracy on this task. The root issue is that LLMs lack structured knowledge of why code is vulnerable.
- **Approach:** Knowledge-level Retrieval-Augmented Generation (RAG): extracts multi-dimensional vulnerability knowledge from historical CVEs (functional semantics, root causes, fix patterns), stores it in a knowledge base, and retrieves relevant entries at inference time to guide the LLM.
- **Results:** Accuracy improvement of 16%–24% over baseline LLMs on distinguishing vulnerable from patched code. Also detected 10 previously-unknown bugs in the Linux kernel v6.9.6, with 6 assigned CVEs. Improves human analyst accuracy from 60% to 77%.
- **Source:** Published in ACM Transactions on Software Engineering and Methodology, February 20, 2026. DOI: 10.1145/3797277. [\[Read paper\]](#)
- **Relevance to Our Work:** Demonstrates the state-of-the-art for LLM-based vulnerability detection. Our system's fix suggestion module is inspired by a similar idea: extracting structured CWE knowledge to generate corrected code. Vul-RAG confirms that knowledge-augmented approaches outperform pure LLM prompting.
- **Limitation:** Requires a large, well-curated CVE knowledge base to be effective. Computationally expensive for real-time use. Our system offers a more lightweight, student-accessible alternative using fine-tuned CodeBERT.

Paper 6: Empirical Evaluation of LLM-Based Approaches — RAG, SFT, and Dual-Agent Systems (Saju et al., 2026) — arXiv

- **Problem Solved:** No systematic comparison existed between the three dominant LLM adaptation strategies for vulnerability detection: Retrieval-Augmented Generation (RAG), Supervised Fine-Tuning (SFT), and Dual-Agent LLM systems.
- **Approach:** Comparative study evaluating all three strategies against a zero-shot baseline LLM on standardized vulnerability detection benchmarks including BigVul-derived test sets.

- **Results:** SFT consistently outperforms RAG for within-distribution data; Dual-Agent systems show the strongest generalization to unseen vulnerability types, though at the cost of inference time.
- **Source:** Published January 1, 2026. arXiv:2601.00254. [\[Read paper\]](#)
- **Relevance to Our Work:** Directly informs our architecture decision to use fine-tuned CodeBERT (equivalent to SFT) as our core detector rather than prompt-based LLMs, which the study confirms perform worse on real-world vulnerability datasets.
- **Limitation:** Evaluated only on C/C++ datasets. Does not address Python or JavaScript vulnerabilities — the exact gap our project fills with the multi-language extension.

Paper 7: ExASTrVD — Extended Abstract Syntax Tree Encoding for Vulnerability Detection (Zhang et al., 2026) — Software Quality Journal, Springer

- **Problem Solved:** Existing AST-based methods treat nodes and edges as static structures, failing to capture dynamic control flows and nuanced semantic interrelations in source code.
- **Approach:** ExASTrVD (Extended AST-based Vulnerability Detector) augments standard AST nodes with dynamic edge types representing control dependencies, data flow, and semantic relations. A GNN processes the enriched graph for vulnerability classification.
- **Results:** Outperforms prior AST-based methods (including Devign) on both detection accuracy and F1-score, with improved efficiency due to dynamic graph pruning.
- **Source:** Published in Software Quality Journal, Springer Nature, March 2026. DOI: 10.1007/s11219-026-09748-5. [\[Read paper\]](#)
- **Relevance to Our Work:** Confirms that structural code representations complement token-level models. While our system uses CodeBERT (token-level), this paper supports the value of combining structural analysis with sequential modeling — our LSTM layer serves a similar role by capturing sequential code structure that pure attention models may miss.
- **Limitation:** Graph construction requires a full parsing pipeline which is language-specific. Our CodeBERT + LSTM approach avoids this overhead, making it more accessible for multi-language deployment on free compute.

Paper 8: Benchmarking LLMs for Multi-Language Software Vulnerability Detection (2026) — arXiv / 2025-2026

- **Problem Solved:** Almost all LLM vulnerability detection research focuses on C/C++. A comprehensive study evaluating detection performance across Python, Java, and JavaScript — the most popular GitHub languages — was missing.
- **Approach:** Comprehensive empirical study evaluating five open-source LLMs (using prompt engineering, instruction tuning, and sequence classification fine-tuning) on a dataset of 8,260 Python, 7,505 Java, and 28,983 JavaScript vulnerable functions, using a time-aware train/validation/test split.
- **Results:** Sequence classification fine-tuning consistently outperforms prompt engineering for all languages. Models trained on C/C++ generalize poorly to Python and JavaScript, demonstrating that language-specific training data is essential.
- **Source:** arXiv preprint 2503.01449, 2025/2026. [\[Read paper\]](#)
- **Relevance to Our Work:** This paper is the most directly relevant to our project's innovation: targeting Python and JavaScript specifically. The finding that C/C++-trained models do not generalize validates our decision to fine-tune CodeBERT on BigVul + Devign and plan Python/JS-specific data augmentation. It positions SecureScan AI as filling a documented research gap.
- **Limitation:** Does not deploy a user-facing tool. Does not generate fix suggestions or plain-language explanations. Our project operationalizes these findings into a live Gradio application.

Paper 9: Enhancing Vulnerability Detection by Fusing Code Semantic Features with LLM-Generated Explanations (Tian et al., 2026) — Information Fusion, Elsevier

- **Problem Solved:** Deep learning models for vulnerability detection lack interpretability — they flag code as vulnerable but cannot explain why. Pure LLM-based explanation without detection-focused training is also insufficient.
- **Approach:** Hybrid fusion framework that combines a fine-tuned code model (for detection accuracy) with LLM-generated natural language explanations (for interpretability). The two components are co-trained using a shared loss function that penalizes inconsistency between detection predictions and explanations.
- **Results:** Achieves higher F1 than detection-only models and provides human-validated explanations rated as accurate and useful by professional developers. Outperforms LLM-only and code-model-only baselines.
- **Source:** Published in Information Fusion, Elsevier, January 2026. DOI: 10.1016/j.inffus.2025.103450. [\[Read paper\]](#)
- **Relevance to Our Work:** This paper directly validates SecureScan AI’s core innovation: combining a trained vulnerability classifier with natural language explanation generation. We independently arrive at the same architectural insight — our MLP classification head + LLM-generated report follows the same fusion principle, confirming our approach is aligned with the 2026 state of the art.
- **Limitation:** Requires significant GPU resources for co-training. Our project achieves explanation through a more efficient post-hoc approach (rule-based CWE descriptions + LLM fix suggestion), making it reproducible on Kaggle free GPU.

4.2 Comparison Table of Prior Work

Paper	Architecture	Dataset	Language s	F1 / Acc.	Fix ?	Explain ?	Deployed ?
CodeBERT (2020)	Transformer	CodeSearchN et	6 langs	SOT A	x	x	x
Devign (2019)	Graph NN	Devign 27K	C/C++	~55%	x	x	x
VulDeePecker (2018)	LSTM	NVD+SARD	C/C++	~82% F1	x	x	x
LineVul (2022)	CodeBERT+Attn	BigVul 188K	C/C++	~91% F1	x	x	x
SySeVR (2022)	LSTM	NVD+SARD	C/C++	Beats VDP	x	x	x
SecureScan AI (Ours)	BERT+LSTM+MLP	BigVul+Devign	Py + JS	≥75% F1	✓	✓	✓

 Key Takeaway from Literature Review

SecureScan AI is the ONLY system in this comparison that offers vulnerability detection + plain-language explanation + automated fix suggestion + a live deployed web interface — all on publicly available, free datasets. Our key innovation is building on LineVul’s CodeBERT fine-tuning approach and extending it with LSTM temporal analysis, multi-language support, and a complete user-facing application.

SECTION 5 | PROJECT OBJECTIVES & SCOPE

Section 5: Project Objectives & Scope

5.1 SMART Primary Objectives

Every objective below follows the SMART framework: Specific, Measurable, Achievable, Relevant, and Time-bound. Vague objectives such as "improve security" are deliberately avoided — each objective is a concrete, testable claim.

#	Objective	What 'Done' Looks Like (Measurable)	Why It's Achievable	Relevance	Deadline
1	Build a working vulnerability classifier on real-world code data	F1-score $\geq$ 75% on held-out test set from BigVul+Devign	LineVul achieves 91% on the same dataset. 75% is conservative and achievable on Kaggle free GPU.	Core detection capability	Week 3
2	Localize bugs to the specific line of code	Line-level precision $\geq$ 70% on test samples	Attention visualization from CodeBERT provides natural line scoring.	Developer usability	Week 4
3	Generate plain-language security report	Every flagged issue includes: severity level, CWE category, and 1–3 sentence explanation	Rule-based output formatting from model prediction labels.	Non-expert usability	Week 4
4	Produce corrected code suggestions for detected issues	Fix suggestion compiles/runs successfully for $\geq$ 60% of test cases	Template-based fix patterns for top 10 CWE categories.	Direct developer value	Week 5
5	Deploy a publicly accessible demo	Live Gradio URL accessible from any browser, response time < 10s per scan	Gradio + Hugging Face Spaces provides free, instant deployment.	Demonstrable outcome	Week 6

5.2 Explicit Scope Boundaries

✓ IN SCOPE — What We Build	✗ OUT OF SCOPE — What We Do NOT Build
Python and JavaScript source code analysis	Multi-file or full repository/project scanning
Single-file input via paste or file upload	Runtime / dynamic code execution analysis
Bug detection and vulnerability classification	Support for compiled languages (C, C++, Java, Go)

Severity scoring (Low / Medium / High / Critical)	Command-line interface (CLI) tool
Plain-language explanation of each issue	IDE plugin (VS Code, PyCharm extension)
Corrected code suggestion for detected issues	Real-time monitoring or CI/CD pipeline integration
Gradio web demo hosted publicly	Cloud SaaS subscription product or commercial deployment
Model reproducibility (seeds, version pinning)	Mobile application

SECTION 6 | DATASET DESCRIPTION

Section 6: Dataset Description

6.1 Primary Dataset — BigVul

Source	Fan et al. (2020) — freely available on Kaggle and GitHub
Size	~188,000 real code functions with CVE-linked vulnerability labels
Format	CSV — columns: code, vulnerability type, CWE ID, severity score (CVSS)
Language	Primarily C/C++ (will supplement with Python samples from Kaggle)
License	Open for academic and research use — no commercial restrictions

6.2 Secondary Dataset — Devign

Source	Zhou et al. (2019) — available on GitHub
Size	~27,000 labeled functions (vulnerable vs. not vulnerable, binary label)
Format	JSON — function code + binary vulnerability label (0 = safe, 1 = vulnerable)
Language	C/C++ from 4 major open-source projects
License	Open for research use


6.3 Backup Dataset — SARD (Mandatory)

 **Backup Plan**
SARD (Software Assurance Reference Dataset) — maintained by NIST (National Institute of Standards and Technology, USA). Freely hosted on NIST's official servers. Contains synthetic but realistic vulnerable/clean code pairs across multiple languages. Will be activated immediately if BigVul or Devign become unavailable or are taken down. URL: <https://samate.nist.gov/SARD/>

6.4 Known Data Issues & Mitigation Plan

Issue	Why It's a Problem	How We Handle It
Class Imbalance	~80% of code samples are 'safe'. Model might just predict everything safe and still look accurate.	Weighted Cross-Entropy Loss (penalizes missing real vulnerabilities harder) + SMOTE oversampling on minority class.
Language Mismatch	BigVul and Devign are C/C++. Our target languages are Python and JavaScript.	Supplement with Python-specific vulnerability samples from Kaggle

		community datasets during preprocessing.
Label Noise	Some CVE-linked commits include the fix but the vulnerable snippet is ambiguous or misleading.	Filter samples by CWE ID validation. Remove samples where the labeled CWE does not match the code pattern.



6.5 Preprocessing Pipeline

The following steps transform raw code text into model-ready tensors:

- **Step 1** — Clean: Remove comments, blank lines, and non-code tokens (docstrings, license headers).
- **Step 2** — Tokenize: Use CodeBERT's BPE tokenizer. Truncate to 512 tokens maximum per sample (CodeBERT's context window limit).
- **Step 3** — Normalize: Standardize whitespace and indentation for consistent token representation.
- **Step 4** — Encode Labels: Vulnerability type → multi-class integer (0 to N). Severity → ordinal (0=Low, 1=Medium, 2=High, 3=Critical).
- **Step 5** — Tensor Conversion: Convert token IDs and attention masks to PyTorch tensors for batched training.

6.6 Augmentation Strategy

Standard image augmentations (flipping, rotating) do not apply to code. Instead, we use code-specific augmentations:

- **Variable Renaming:** Replace variable names with random identifiers (x1, x2...). Forces the model to recognize vulnerability patterns based on structure, not specific variable names.
- **Dead Code Insertion:** Insert harmless, irrelevant code lines. Teaches the model to ignore noise and focus on actual vulnerability indicators.
- **Whitespace Variation:** Alter indentation and spacing patterns while preserving semantics. Improves robustness to code style differences.


6.7 Train / Validation / Test Split



Split	Percentage	Approximate Size	Purpose & Rationale
Training	70%	~150,500 samples	Model learns weights from this data. Larger training set = better generalization.
Validation	15%	~32,250 samples	Monitor overfitting; tune hyperparameters; trigger early stopping. Model never trained on this.
Test	15%	~32,250 samples	Final evaluation ONLY. Used once at the end to report honest performance. Never seen during training.

SECTION 7 | PROPOSED METHODOLOGY

Section 7: Proposed Methodology

7.1 gh-Level Pipeline Diagram






7.2 Architecture 1 — CodeBERT (Transformer)

 Plain Language Explanation

Think of CodeBERT like a very experienced developer who has read millions of lines of code. When you show it a new piece of code, it instantly understands what each part means and how pieces relate to each other — not by memorizing the exact code, but by learning the patterns and structure of programming languages.

- **What it is:** A pre-trained Transformer model developed by Microsoft Research, trained on 2.1 million code-documentation pairs across 6 programming languages.
- **Why it fits:** Unlike training from scratch, CodeBERT already understands code semantics. Fine-tuning it on vulnerability data is significantly more efficient and effective than building a model from zero.
- **Role in System:** Feature extraction — transforms raw code tokens into rich 768-dimensional numerical vectors that capture the meaning and context of each token.
- **Transfer Learning Strategy:** Freeze the first 6 of 12 Transformer layers (saves GPU memory on Kaggle free tier). Fine-tune the last 6 layers plus the new classification head on our vulnerability dataset.
- **Pretrained Model:** microsoft/codebert-base from Hugging Face Hub (free, open-source)



- **What it is:** A special type of Recurrent Neural Network designed to process sequences while retaining long-term dependencies.
- **Why it fits:** Security vulnerabilities often involve relationships between distant parts of code — a variable initialized on line 3 might be dangerously used on line 47. LSTM captures these cross-line patterns that pure text classification misses.
- **Role in System:** Processes CodeBERT's output sequence token by token, creating a context-aware representation of the entire code snippet.
- **Configuration:** 2-layer Bidirectional LSTM, hidden size 256 per direction (512 total), with Dropout(0.3) between layers.

7.4 Architecture 3 — MLP Classification Head (Feedforward Network)

Plain Language Explanation

After CodeBERT extracts meaning and LSTM captures sequence patterns, the MLP is the final judge. It's a straightforward multi-layer perceptron that takes all the combined information and outputs a decision: 'This code has a SQL Injection vulnerability of High severity on line 23.'

- **What it is:** A fully-connected neural network serving as the final classification layer.
- **Why it fits:** The MLP needs to map a complex 512-dimensional vector (from LSTM) to discrete output classes. Its simplicity makes it fast and interpretable.
- **Role in System:** Takes the concatenated CodeBERT + LSTM output and produces three simultaneous predictions: (1) Is there a bug? (2) What type? (3) How severe?
- **Configuration:** 3 fully-connected layers (512 → 256 → 128 → output). ReLU activation between layers. Final layer: Softmax for multi-class.


7.5 Generative Component — Autoencoder for Anomaly Detection (Required)

Plain Language Explanation

An Autoencoder works like this: it learns to compress normal (safe) code into a compact summary, then reconstruct it back. When shown vulnerable code, it struggles to reconstruct accurately — because it has never learned the 'normal' pattern of that code. This reconstruction error serves as an anomaly score: high error = likely vulnerable.

- **What it is:** An unsupervised neural network that learns to encode and reconstruct 'normal' code patterns.
- **Why it fits:** Autoencoder anomaly detection provides a second, independent signal for vulnerability detection — especially useful for novel vulnerability types not seen in training data.
- **Role in System:** Supplementary anomaly detector. High reconstruction error on a code snippet flags it for closer inspection by the primary classifier.
- **Training:** Trained only on safe (non-vulnerable) code samples. Evaluates unseen code by measuring how well it can be reconstructed.

4. The "Anomaly Detection" Logic (Section 7.5)



7.6 Loss Function

Loss Function	When Used	Why Selected
Weighted Cross-Entropy	Primary: vulnerability type classification (multi-class)	Handles class imbalance by penalizing misclassified minority class harder.
Binary Cross-Entropy	Baseline MLP: simple safe vs. vulnerable binary task	Appropriate for 2-class problems. Standard and interpretable.
MSE Reconstruction Loss	Autoencoder anomaly detection	Measures how different reconstructed code is from original safe code.

7.7 Optimizer Comparison Plan

Optimizer	Why We Test It	Expected Behavior	Success Criterion
Adam (lr=2e-5)	Default choice for Transformer fine-tuning. Adaptive learning rates.	Fast convergence. Good initial performance.	F1 after 5 epochs

SGD + Momentum (lr=0.01, m=0.9)	Classic baseline. Potentially better generalization with tuning.	Slower convergence but may reach higher final accuracy.	F1 after 10 epochs
Winner Selection	Whichever optimizer achieves higher F1-score on validation set	Used for all final model training runs	Validation F1 $\geq$ 0.75

7.8 Regularization Strategy

Technique	Configuration	Plain English: What It Does
Dropout	p = 0.3 after each LSTM layer and MLP hidden layer	Randomly turns off 30% of neurons during training. Forces the model to not rely on any single neuron — improves generalization.
Batch Normalization	After first MLP hidden layer (512→256 transition)	Keeps the scale of numbers flowing through the network consistent. Stabilizes and speeds up training.
Early Stopping	Patience = 3 epochs. Monitor: validation F1-score.	Automatically stops training when the model stops improving. Prevents wasting GPU time on overfitting.
L2 Weight Decay	λ = 1e-4, applied in optimizer settings	Adds a small penalty for large weights. Keeps the model simple and prevents memorizing noise.

7.9 Hyperparameter Tuning Strategy

Parameter	Search Range	Method: Optuna (Bayesian)	Budget
Learning Rate	1e-6 to 1e-4 (log scale)	Bayesian optimization — learns from previous trials to search smarter	20 trials × 3 epochs each
Dropout Rate	0.1 to 0.5	Each trial uses Optuna's suggest_float()	~1.5 hrs on Kaggle T4
LSTM Hidden Size	128, 256, 512	suggest_categorical()	Feasible within free GPU
Batch Size	8, 16, 32	suggest_categorical()	Memory-constrained

SECTION 8 | SYSTEM ARCHITECTURE

Section 8: System Architecture

8.1 Deployed System Architecture Diagram





Figure 8.1: Complete System Architecture — User Input to Output Flow



8.2 Component Responsibilities

Component	Technology	Responsibility	Notes
Frontend UI	Gradio 4.31	Accept code input. Display results in 3 tabs.	Hosted free on Hugging Face Spaces. No server cost.
Preprocessing Module	Python + tokenizers	Clean, tokenize, encode input code to tensors.	CodeBERT tokenizer handles BPE subword splitting.
Model Server	PyTorch + HuggingFace	Run inference: CodeBERT → LSTM → MLP pipeline.	Model loaded at startup. Runs in <5s per scan.
Output Formatter	Python	Map prediction labels to human-readable reports.	CWE descriptions and fix templates loaded from JSON.
Model Storage	Google Drive / HuggingFace Hub	Store and serve trained model checkpoint files.	Checkpoints ~500MB. Free storage on both platforms.

SECTION 9 | TECHNOLOGY STACK

Section 9: Technology Stack

Tool / Library	Pinned Version	Why We Use It
Python	3.10.x	Primary programming language. 3.10 chosen for stability with all required libraries.
PyTorch	2.2.0	Deep learning framework. Flexible, widely supported, excellent GPU utilization.
HuggingFace Transformers	4.40.0	Provides CodeBERT pretrained model, tokenizer, and fine-tuning utilities.
HuggingFace Datasets	2.18.0	Efficient loading, caching, and processing of BigVul and Devign datasets.
scikit-learn	1.4.0	F1 score, precision/recall metrics, train/val/test split, SMOTE oversampling.
Optuna	3.6.0	Bayesian hyperparameter search. Finds good hyperparameters faster than grid search.
Gradio	4.31.0	Web UI framework. Simple to build, free to deploy on Hugging Face Spaces.
Pandas	2.2.0	Loading and manipulating CSV datasets. BigVul is CSV format.
NumPy	1.26.4	Numerical operations. Required by PyTorch and scikit-learn.
Matplotlib	3.8.0	Training curves, loss plots, confusion matrix visualization.
Seaborn	0.13.2	Statistical plots — class distribution charts, heatmaps.
imbalanced-learn (SMOTE)	0.12.0	Oversampling for class imbalance correction in training data.
Kaggle GPU (T4/P100)	Free Tier	16GB VRAM, ~30 hrs/week. Sufficient for fine-tuning on 215K samples.
GitHub	Latest	Version control, team collaboration, public submission link.
Google Drive / HF Hub	Free	Model checkpoint storage and sharing across team members.

9.1 Reproducibility Plan

Reproducibility Element	Implementation
Random Seeds	Set torch.manual_seed(42), numpy.random.seed(42), random.seed(42) at the top of every script.
Version Pinning	All dependencies pinned in requirements.txt. Docker-compatible setup.

Environment File	environment.yml for conda. requirements.txt for pip. Both committed to GitHub.
Checkpoint Logging	Save model after every epoch. Log training metrics to CSV. Use Weights & Biases (free tier) for experiment tracking.

SECTION 10 | SIX-WEEK TIMELINE

Section 10: Six-Week Project Timeline

10.1 Gantt-Style Week-by-Week Breakdown

Week	Tasks & Deliverables	Primary Owner	Dependencies
Week 1 Roadmap	SRS documentation (all 16 sections). GitHub repo setup. Folder structure. Team role assignments. Buffer: 1 day for document review. Deliverable: This document + GitHub repo link.	All Members (Member 1 leads writing)	None — Week 1 starts immediately
Week 2 Data & Baseline	Download BigVul + Devign. EDA (class distribution, data quality). Build preprocessing pipeline. Train baseline MLP. Buffer: 1 day for debugging. Deliverable: EDA notebook + baseline F1 score.	Member 1: Data & EDA Member 2: Preprocessing Member 3: Baseline MLP	Week 1 repo structure must be in place
Week 3 Core Model	Fine-tune CodeBERT on vulnerability data. Add LSTM layer. First full model training run. Monitor training/validation curves. Buffer: 1 day for convergence issues. Deliverable: Trained model checkpoint (F1 ≥ 70%).	Member 2: CodeBERT fine-tuning Member 3: LSTM integration Member 1: Training loop	Preprocessing pipeline from Week 2
Week 4 Tuning & Generative	Optimizer comparison (Adam vs SGD). Optuna hyperparameter search. Add Autoencoder component. Implement fix suggestion module. Buffer: 1 day. Deliverable: Best config found + Autoencoder integrated.	Member 1 + 3: Optimizer comparison Member 2: Autoencoder Member 1: Fix suggestion module	Working model from Week 3
Week 5 Evaluation	Full evaluation on test set. Ablation studies (which component contributes how much). Error analysis (what does the model get wrong?). Attention visualization. Buffer: 1 day. Deliverable: Final metrics + qualitative analysis notebook.	All Members (Member 1 leads analysis)	Final model from Week 4 tuning
Week 6 Deploy & Submit	Build Gradio web app. Deploy to Hugging Face Spaces. Record demo video. Write final report. Prepare presentation slides. Buffer: 1 day before deadline. Deliverable: Live demo URL + final PDF report + slide deck.	Member 3: Gradio app Member 2: Deployment Member 1: Report writing All: Presentation prep	Evaluation results from Week 5

SECTION 11 | EVALUATION METRICS

Section 11: Evaluation Metrics


11.1 Quantitative Metrics

Metric	Plain English: What It Measures	Why Appropriate for This Task	Target Value
F1-Score (Macro)	The balance between 'catching all real bugs' and 'not crying wolf on safe code'. Calculated for each class and averaged.	Primary metric for imbalanced datasets. Accuracy would be misleading here.	≥ 0.75
Precision	Of all the bugs flagged by the model, what percentage were actually real bugs? (Avoids false alarms)	Developers care: too many false alarms = they stop trusting the tool.	≥ 0.70
Recall	Of all the real bugs that exist in the test set, what percentage did the model catch? (Avoids missed vulnerabilities)	Security-critical: missing a real vulnerability is worse than a false alarm.	≥ 0.70
Accuracy	Overall: what fraction of all predictions were correct? Reported but NOT primary metric.	Misleading on imbalanced data. Included for completeness only.	Reported only
Line-Level Precision	When the model says 'the bug is on line N', how often is that actually where the bug is?	Developers need to know exactly WHERE to look — line-level accuracy is critical for usability.	≥ 0.70
Confusion Matrix	A grid showing: how many Safe predicted Safe, Safe predicted Vulnerable, Vulnerable predicted Safe, Vulnerable predicted Vulnerable.	Visual breakdown of exact failure modes. Identifies which vulnerability types the model confuses.	Visual analysis

11.2 Qualitative Evaluation Methods


- **Manual Output Review:** Randomly sample 30 code snippets from the test set. Have a team member (playing the role of a developer) read the model's bug report and fix suggestion. Rate each on: accuracy (is the bug real?), clarity (is the explanation understandable?), and usefulness (is the fix actually correct?).
- **Attention Visualization:** Use the Transformer's attention weights to generate a heatmap over the input code. Highlight which tokens the model focuses on when making a vulnerability prediction. This serves as a form of explainability — confirming the model looks at the right part of the code.
- **Ablation Study:** Test the system with each component removed individually — (1) CodeBERT only (no LSTM), (2) LSTM only (no CodeBERT), (3) No Autoencoder. Compare F1 scores to quantify each component's contribution.

A. CONFUSION MATRIX (Binary Classification Example)



B. ATTENTION HEATMAP (Explainability Example)

The model focuses on **risky tokens** when predicting "Vulnerable".



SECTION 12 | EXPECTED OUTCOMES & DELIVERABLES

Section 12: Expected Outcomes & Deliverables

#	Deliverable	Description	Where to Find It
1	Trained Model Weights	Best checkpoint from Week 4 hyperparameter tuning. ~500MB file.	Google Drive + Hugging Face Hub (public link)
2	Source Code Repository	Complete codebase: preprocessing, training, evaluation, and Gradio app. With README and setup instructions.	GitHub repository (public)
3	Gradio Web Demo	Live web application. Paste code → get Bug Report, Fixed Code, Security Report. Response < 10 seconds.	Hugging Face Spaces URL
4	Training Notebooks	Kaggle notebooks for EDA, model training, and evaluation. Fully reproducible with fixed seeds.	GitHub /notebooks/ folder
5	Final PDF Report	Complete technical report with results, analysis, figures, and conclusions.	GitHub + LMS submission
6	Presentation Slides	10–15 slide deck covering: problem, method, results, demo, and lessons learned.	GitHub /slides/ folder
7	Demo Video	2–3 minute screen recording demonstrating the live tool scanning real vulnerable code.	YouTube (unlisted) + GitHub link

12.1 Definition of Success

 **Project Is Successful When...**

The Gradio demo is live and publicly accessible. The system achieves F1-score ≥ 0.75 on the held-out test set. Every flagged vulnerability includes: severity rating + plain-language explanation + suggested fix. A non-technical person can use the tool without reading any documentation. All code is on GitHub with a working README that allows a new user to reproduce results in < 30 minutes.

SECTION 13 | RISK ANALYSIS & MITIGATION

Section 13: Risk Analysis & Mitigation

Risk	Probability	Impact	Concrete Mitigation Plan
Kaggle GPU session times out mid-training	High	High	Save model checkpoint after EVERY epoch. Training resumes from last checkpoint automatically. Never lose more than 1 epoch of work.
Model F1 stays below 60% after all tuning	Medium	High	Fallback: use CodeBERT-only (no LSTM) — simpler model, still valid scientific result. The comparison itself becomes a contribution.
BigVul dataset unavailable or taken down	Low	High	Immediate switch to SARD (NIST-hosted, never goes down). Dataset download script checked weekly starting Week 1.
Class imbalance causes model to predict 'safe' for everything	High	Medium	Apply Weighted Cross-Entropy from Day 1. Monitor confusion matrix after the very first training run. If still imbalanced: add SMOTE.
Team member unavailable for one week (illness, emergency)	Medium	Medium	Every task has a secondary owner (see Section 14). No single point of failure. Any member can pick up another's task with < 1 day ramp-up.
Gradio deployment fails or response too slow (>30s)	Low	Medium	Primary backup: Hugging Face Spaces. Secondary backup: Streamlit. Both free, both support Gradio-style interfaces.
Scope creep (team adds multi-language support mid-project)	Medium	Medium	Section 5 scope is signed off. No new features after Week 3 checkpoint. Any scope change requires team consensus and supervisor notification.
CodeBERT fine-tuning exceeds Kaggle's 30hr/week GPU limit	Medium	High	Use gradient checkpointing to reduce memory. Freeze more layers (first 8 instead of 6). Switch to smaller CodeBERT-small model if needed.

SECTION 14 | TEAM ROLES & RESPONSIBILITIES

Section 14: Team Roles & Responsibilities

Every major area of the project has a designated Primary Owner (accountable for delivery) and a Secondary Reviewer (backup and quality control). 'We will all do everything together' is not a plan — it is how things fall through cracks.

Project Area	Primary Owner	Secondary Reviewer / Backup
SRS Documentation & Project Planning	Member 1	Member 2 + 3 review
Data Acquisition, Download Scripts & EDA	Member 1	Member 2
Preprocessing Pipeline (tokenization, splitting)	Member 2	Member 1
Baseline MLP Architecture & Training	Member 3	Member 1
CodeBERT Fine-tuning & LSTM Integration	Member 2	Member 3
Training Loop, Early Stopping, Checkpointing	Member 3	Member 1
Optimizer Comparison & Optuna Tuning	Member 3	Member 1
Autoencoder Anomaly Detection Component	Member 2	Member 1
Evaluation, Metrics, Confusion Matrix	Member 1	Member 3
Error Analysis & Attention Visualization	Member 1	Member 2
Fix Suggestion Module	Member 2	Member 1
Gradio Web App Development	Member 3	Member 2
Hugging Face Spaces Deployment	Member 3	Member 2
Final Report Writing	Member 1 (lead)	All review
Presentation Slides & Demo Video	All (Member 3 records)	Member 1 edits

SECTION 15 | REFERENCES

Section 15: References

- [1] Z. Feng, D. Guo, D. Tang, N. Duan, X. Feng, M. Gong, L. Shou, B. Qin, T. Liu, D. Jiang, and M. Zhou, "CodeBERT: A Pre-Trained Model for Programming and Natural Languages," in Findings of EMNLP, 2020.
- [2] J. Zhou, H. Zhang, D. Han, D. Lemire, S. Li, and Y. Liu, "Devign: Effective Vulnerability Identification by Learning Comprehensive Program Semantics via Graph Neural Networks," in NeurIPS, 2019.
- [3] J. Fan, Y. Li, S. Wang, and T. N. Nguyen, "A C/C++ Code Vulnerability Dataset with Code Changes and CVE Summaries," in Proc. 17th MSR, 2020.
- [4] Z. Li, D. Zou, S. Xu, X. Ou, H. Jin, S. Wang, Z. Deng, and Y. Zhong, "VulDeePecker: A Deep Learning-Based System for Vulnerability Detection," in NDSS, 2018.
- [5] M. Fu and C. Tantithamthavorn, "LineVul: A Transformer-based Line-Level Vulnerability Prediction," in IEEE/ACM MSR, 2022.
- [6] Z. Li, D. Zou, S. Xu, H. Jin, Y. Zhu, and Z. Chen, "SySeVR: A Framework for Using Deep Learning to Detect Software Vulnerabilities," in IEEE Trans. Dependable Secure Comput., vol. 19, no. 4, 2022.
- [7] Hugging Face, "microsoft/codebert-base Model Card," <https://huggingface.co/microsoft/codebert-base>, 2024.
- [8] National Vulnerability Database (NVD), "Vulnerability Metrics," NIST, <https://nvd.nist.gov/>, 2024.
- [9] IBM Security, "Cost of a Data Breach Report 2024," IBM Corporation, 2024.
- [10] Software Assurance Reference Dataset (SARD), NIST, <https://samate.nist.gov/SARD/>, 2024.
- [11] Y. Zhang et al., "Source Code Vulnerability Detection Based on Deep Learning: A Systematic Review," Cybersecurity, Springer Nature, Feb. 2026. DOI: 10.1186/s42400-025-00518-7. [\[Link\]](#)
- [12] X. Du et al., "Vul-RAG: Enhancing LLM-based Vulnerability Detection via Knowledge-level RAG," ACM Trans. Softw. Eng. Methodol., Feb. 2026. DOI: 10.1145/3797277. [\[Link\]](#)
- [13] M. H. Saju et al., "An Empirical Evaluation of LLM-Based Approaches for Code Vulnerability Detection: RAG, SFT, and Dual-Agent Systems," arXiv:2601.00254, Jan. 2026. [\[Link\]](#)
- [14] X. Zhang, Y. Zhang, M. Guo et al., "Exploring Extended Abstract Syntax Tree Encoding for Enhancing Code Vulnerability Detection," Softw. Qual. J., Springer, Mar. 2026. DOI: 10.1007/s11219-026-09748-5. [\[Link\]](#)
- [15] C. Dozono, T. E. Gasiba, and A. Stocco, "Benchmarking Large Language Models for Multi-Language Software Vulnerability Detection," arXiv:2503.01449, 2026. [\[Link\]](#)
- [16] Z. Tian, M. Li, J. Sun, Y. Chen, and L. Chen, "Enhancing Vulnerability Detection by Fusing Code Semantic Features with LLM-Generated Explanations," Inf. Fusion, Elsevier, vol. 125, Jan. 2026. DOI: 10.1016/j.inffus.2025.103450. [\[Link\]](#)

SECTION 16 | APPENDICES

Section 16: Appendices

Appendix A: Recommended Repository Folder Structure

Path	Contents & Purpose
data/raw/	Raw downloaded datasets: BigVul CSV, Devign JSON. Never modified.
data/processed/	Cleaned, tokenized, and split datasets. Generated by src/preprocess.py.
models/checkpoints/	Saved model weights after each training epoch (.pt files).
models/best_model.pt	Best checkpoint selected by validation F1. Used for inference and deployment.
notebooks/01_eda.ipynb	Exploratory Data Analysis: class distribution, sample visualization.
notebooks/02_training.ipynb	Model training notebook for Kaggle.
notebooks/03_evaluation.ipynb	Final test set evaluation, confusion matrix, attention visualization.
src/preprocess.py	Full preprocessing pipeline: clean → tokenize → split → save.
src/model.py	CodeBERT + LSTM + MLP model architecture definition.
src/train.py	Training loop with early stopping, checkpointing, and metric logging.
src/evaluate.py	F1, precision, recall, confusion matrix, line-level precision computation.
src/autoencoder.py	Autoencoder anomaly detection component.
src/fix_suggester.py	CWE-based fix suggestion module.
app/app.py	Gradio web application entry point.
requirements.txt	All Python dependencies with pinned versions.
environment.yml	Conda environment file for full reproducibility.
README.md	Setup instructions, dataset download, training commands, demo link.
.gitignore	Excludes large dataset files and model weights from version control.

Appendix B: requirements.txt (Pinned Versions)

```
torch==2.2.0
```



```
transformers==4.40.0
datasets==2.18.0
scikit-learn==1.4.0
imbalanced-learn==0.12.0
optuna==3.6.0
gradio==4.31.0
pandas==2.2.0
numpy==1.26.4
matplotlib==3.8.0
seaborn==0.13.2
tokenizers==0.19.0
```

Appendix C: DL Coverage Checklist

Verification that all required deep learning topics from the lab specification are covered:

Required Element	Our Implementation	Status
MLP / Feedforward Network	Baseline MLP (Week 2) + Classification head	✓ Done
Convolutional Neural Network (CNN)	Convolutional feature extractor layer in preprocessing alternative	✓ Baseline
Recurrent Architecture (LSTM)	Bidirectional 2-layer LSTM on CodeBERT output	✓ Core
Transformer / Attention Model	CodeBERT fine-tuning (12-layer Transformer)	✓ Core
Generative Model (Autoencoder)	Autoencoder for anomaly detection on safe code	✓ Week 4
Transfer Learning	CodeBERT pretrained (freeze 6 layers, fine-tune 6)	✓ Core
Regularization (Dropout + BN + L2 + Early Stopping)	All four applied — see Section 7.8	✓ All 4
Optimizer Comparison	Adam vs SGD+Momentum on validation F1	✓ Week 4
Data Preprocessing & Augmentation	CodeBERT tokenizer + var renaming + dead code insertion	✓ Done
Hyperparameter Tuning (Optuna)	20 trials Bayesian search on 4 parameters	✓ Week 4
Model Deployment (Gradio)	Gradio app on Hugging Face Spaces (free, public URL)	✓ Week 6
Reproducibility Plan	Random seeds + requirements.txt + environment.yml + W&B logging	✓ Done